MonkDB PDF Vector Ingestion Test Document

This sample PDF is created to test the PDF vector ingestion pipeline.

It contains a few paragraphs of text so that we can split it into chunks and ingest them into MonkDB.

MonkDB enables multi-model analytical workloads, including vector search, without requiring ETL.

The goal of this test is to verify ingestion, embedding generation, and semantic search.